


# Week 1

REFER TO THE GITHUB FOR THE LAB CODES - [https://github.com/Bhardwaj-Saurabh/Machine\\_Learning\\_Specialization\\_AndrewNG\\_Coursera](https://github.com/Bhardwaj-Saurabh/Machine_Learning_Specialization_AndrewNG_Coursera)


## Supervised learning part 2

Two or more inputs



2 or more inputs can be given to a classification supervised model and here the machine would try to fit a line between the 2 classes/categories (here). the learning algorithm will learn to fit a boundary line to this data.

## Unsupervised learning part 1



## Unsupervised learning part 2

### Unsupervised learning

Data only comes with inputs  $x$ , but not output labels  $y$ .  
Algorithm has to find **structure** in the data.

#### Clustering

Group similar data points together.

#### Dimensionality reduction

Compress data using fewer numbers.

#### Anomaly detection

Find unusual data points.

## Linear regression model part 1

## Terminology

Training Data used to train the model

set:	x	y
	size in feet <sup>2</sup>	price in \$1000's
(1)	2104	400
(2)	1416	232
(3)	1534	315
(4)	852	178
...	...	...
(47)	3210	870

$m = 47$

$x^{(1)} = 2104$   $y^{(1)} = 400$   
 $(x^{(1)}, y^{(1)}) = (2104, 400)$   
 $x^{(2)} = 1416$   $x^{(2)} \neq x^2$  not exponent

Notation:

$x$  = "input" variable  
feature

$y$  = "output" variable  
"target" variable

$m$  = number of training examples

$(x, y)$  = single training example

$(x^{(i)}, y^{(i)})$

$(x^{(i)}, y^{(i)})$  =  $i^{\text{th}}$  training example  
(1<sup>st</sup>, 2<sup>nd</sup>, 3<sup>rd</sup> ...)

## C1\_W1\_Lab03\_Model\_Representation\_Soln.ipynb

```
w = 100
b = 100
print(f"w: {w}")
print(f"b: {b}")

def compute_model_output(x, w, b):

    m = x.shape[0] #----> x.shape returns (n,0) where n is the length of the array
    f_wb = np.zeros(m)
    for i in range(m):
        f_wb[i] = w * x[i] + b

    return f_wb
#f_wb is nothing but y values!!! which we plot
```

here we are finding the value of  $f_{wb}$  for every value of  $x$  given a particular value we defined ( $w, b$ ). we then took the size of  $x$ , made an array of that size and then changed the values to get a list of  $f_{wb}$  values. which we plot.


```
tmp_f_wb = compute_model_output(x_train, w, b,)

# Plot our model prediction
plt.plot(x_train, tmp_f_wb, c='b', label='Our Prediction')

# Plot the data points
plt.scatter(x_train, y_train, marker='x', c='r', label='Actual Values')

# Set the title
plt.title("Housing Prices")
# Set the y-axis label
plt.ylabel('Price (in 1000s of dollars)')
# Set the x-axis label
plt.xlabel('Size (1000 sqft)')
plt.legend()
plt.show()
```

## Cost function formula




## Cost function intuition

goal of linear regression:

$$\underset{w}{\text{minimize}} J(w)$$

general case:


$$\underset{w,b}{\text{minimize}} J(w,b)$$



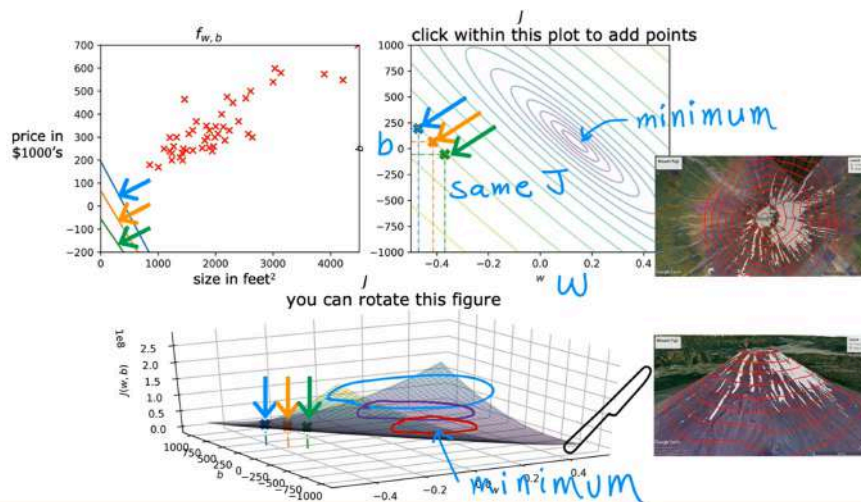
model function is the function of  $X$  for a fixed value of  $w, b$

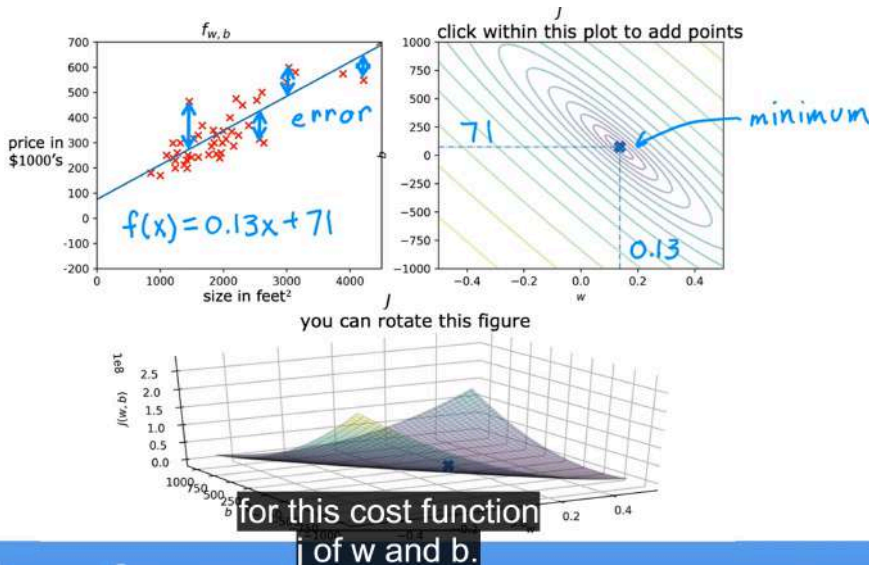
cost function is a function of  $w, b$ ! not  $X$ !

## Cost function visualization

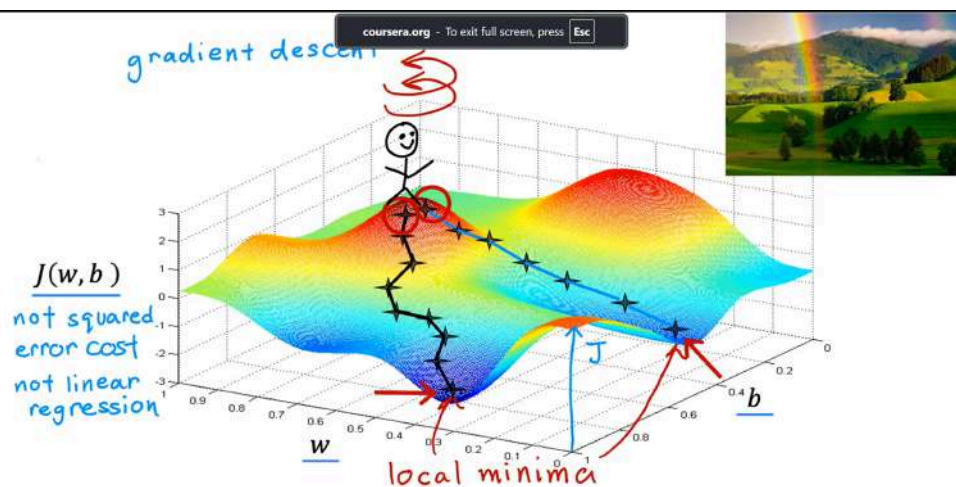


when we plot  $j(w, b)$  we get a 3d plot. when we plot  $j(w)$  keeping  $b=0$ , we get a 2d parabolic graph.





## Gradient descent



here we have an example of a cost fn not "squared error" one but a diff one and also not a LR model (which would give you a parabolic curve to the cost fn. so here too we use GD to get to the local minima of the graph.

the baby steps here is the alpha is the learning rate, how fast you move down the hill.

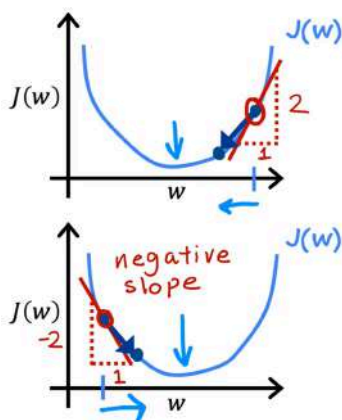
correct and incorrect way to implement gradient descent simultaneous updation of parameters

### Correct: Simultaneous update

$$\begin{aligned} \text{tmp\_w} &= w - \alpha \frac{\partial}{\partial w} J(w, b) \\ \text{tmp\_b} &= b - \alpha \frac{\partial}{\partial b} J(w, b) \\ w &= \text{tmp\_w} \\ b &= \text{tmp\_b} \end{aligned}$$

### Incorrect

$$\begin{aligned} \text{tmp\_w} &= w - \alpha \frac{\partial}{\partial w} J(w, b) \\ w &= \text{tmp\_w} \\ \text{tmp\_b} &= b - \alpha \frac{\partial}{\partial b} J(w, b) \\ b &= \text{tmp\_b} \end{aligned}$$



$$w = w - \alpha \frac{d}{dw} J(w)$$

$> 0$

$$w = w - \alpha \cdot (\text{positive number})$$


$$\frac{d}{dw} J(w)$$

$< 0$

$$w = w - \alpha \cdot (\text{negative number})$$

understanding the derivative used in GD






Can reach local minimum with fixed learning rate  $\alpha$

$w = w - \underbrace{\alpha \frac{d}{dw} J(w)}_{\text{smaller, not as large, large}}$

Near a local minimum,

- Derivative becomes smaller
- Update steps become smaller

Can reach minimum without decreasing learning rate  $\alpha$



## Gradient descent for linear regression

(Optional)

$$\frac{\partial}{\partial w} J(w, b) = \frac{\partial}{\partial w} \frac{1}{2m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})^2 = \frac{\partial}{\partial w} \frac{1}{2m} \sum_{i=1}^m (wx^{(i)} + b - y^{(i)})^2$$


$$= \frac{1}{2m} \sum_{i=1}^m (wx^{(i)} + b - y^{(i)}) \cdot 2x^{(i)} = \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})x^{(i)}$$

$$\frac{\partial}{\partial b} J(w, b) = \frac{\partial}{\partial b} \frac{1}{2m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})^2 = \frac{\partial}{\partial b} \frac{1}{2m} \sum_{i=1}^m (wx^{(i)} + b - y^{(i)})^2$$

$$= \frac{1}{2m} \sum_{i=1}^m (wx^{(i)} + b - y^{(i)}) \cdot 2 = \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})$$

no  $x^{(i)}$

squared error cost



## C1\_W1\_Lab04\_Gradient\_Descent\_Soln

code to implement GD for LR

```

import math, copy
import numpy as np
import matplotlib.pyplot as plt

# Load our data set
x_train = np.array([1.0, 2.0]) #features
y_train = np.array([300.0, 500.0]) #target value

#Function to calculate the cost
def compute_cost(x, y, w, b):

    m = x.shape[0]
    cost = 0

    for i in range(m):
        f_wb = w * x[i] + b
        cost = cost + (f_wb - y[i])**2
    total_cost = 1 / (2 * m) * cost

    return total_cost

"""
You will implement gradient descent algorithm for one feature. You will need three functions.

- 'compute_gradient' implementing equation (4) and (5) above
- 'compute_cost' implementing equation (2) above (code from previous lab)
- 'gradient_descent', utilizing compute_gradient and compute_cost

```

Conventions:

```

- The naming of python variables containing partial derivatives follows this pattern,  $\partial J(w,b) / \partial b$  will be 'dj_db'.
- w.r.t is With Respect To, as in partial derivative of  $J(w,b)$  With Respect To  $b$ .
"""

```

```

def compute_gradient(x, y, w, b):
    """
    Computes the gradient for linear regression
    Args:
        x (ndarray (m,)): Data, m examples
        y (ndarray (m,)): target values
        w,b (scalar)    : model parameters
    Returns
        dj_dw (scalar): The gradient of the cost w.r.t. the parameters w
        dj_db (scalar): The gradient of the cost w.r.t. the parameter b
    """

    # Number of training examples
    m = x.shape[0]
    dj_dw = 0
    dj_db = 0

    for i in range(m):
        f_wb = w * x[i] + b
        dj_dw_i = (f_wb - y[i]) * x[i]
        dj_db_i = f_wb - y[i]
        dj_db += dj_db_i
        dj_dw += dj_dw_i
    dj_dw = dj_dw / m
    dj_db = dj_db / m

    return dj_dw, dj_db

def gradient_descent(x, y, w_in, b_in, alpha, num_iters, cost_function, gradient_function):
    # An array to store cost J and w's at each iteration primarily for graphing later
    #J_history = []
    #p_history = []
    b = b_in
    w = w_in

    for i in range(num_iters):
        # Calculate the gradient and update the parameters using gradient_function
        dj_dw, dj_db = gradient_function(x, y, w, b)

        # Update Parameters using equation (3) above
        b = b - alpha * dj_db
        w = w - alpha * dj_dw

        # Save cost J at each iteration

```

```

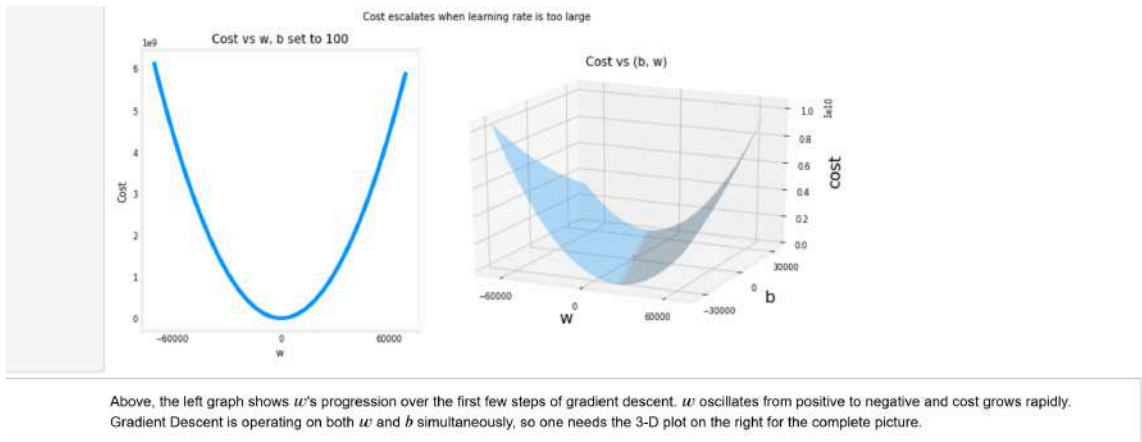
    if i<100000: # prevent resource exhaustion
        J_history.append( cost_function(x, y, w , b))
        #p_history.append([w,b])
    # Print cost every at intervals 10 times or as many iterations if < 10
    """if i% math.ceil(num_iters/10) == 0:
        print(f"Iteration {i:4}: Cost {J_history[-1]:0.2e} ",
              f"dj_dw: {dj_dw: 0.3e}, dj_db: {dj_db: 0.3e} ",
              f"w: {w: 0.3e}, b:{b: 0.5e}")"""

    return w, b, J_history, p_history #return w and J,w history for graphing

#-----client code

# initialize parameters
w_init = 0
b_init = 0
# some gradient descent settings
iterations = 10000
tmp_alpha = 1.0e-2
# run gradient descent
w_final, b_final, J_hist, p_hist = gradient_descent(x_train ,y_train, w_init, b_init, tmp_alpha,
                                                    iterations, compute_cost, compute_gradient)
print(f"(w,b) found by gradient descent: ({w_final:8.4f},{b_final:8.4f})")

```



final code for GD for LR in univariate

```

import numpy as np

x_train = np.array([1.0, 2.0])
y_train = np.array([300.0, 500.0])

w_init = 0
b_init = 0
alpha = 0.001
num_iters = 100000

def costfn(x_train, y_train, w,b):

    j_wb = 0

    m = x_train.shape[0]

    for i in range(m):

        f_wb = w*x_train[i]+b

        j_wb = j_wb + (f_wb-y_train[i])**2

    j_wb = j_wb/(2*m)

```

```

    return j_wb

def gradientfn(x_train, y_train, w,b):

    m = x_train.shape[0]

    d_jwb_w = 0

    d_jwb_b = 0


    for i in range(m):

        f_wb = w*x_train[i] + b

        d_jwb_w = d_jwb_w + (f_wb-y_train[i])*x_train[i]

        d_jwb_b = d_jwb_b + (f_wb-y_train[i])

    d_jwb_w = d_jwb_w/m
    d_jwb_b = d_jwb_b/m


    return d_jwb_w, d_jwb_b


def gradient_descent(x_train, y_train, alpha, w_init, b_init, num_iters, costfn, gradientfn):

    w,b = w_init, b_init

    dj_dw , dj_db = gradientfn(x_train, y_train, w,b)

    jwb = []


    for i in range(num_iters):

        dj_dw , dj_db = gradientfn(x_train, y_train, w,b)

        w = w - alpha*dj_dw

        b = b - alpha*dj_db

        if (i%1000)==0:

            jwb.append(costfn(x_train, y_train, w,b))

    return w,b,jwb


w_final, b_final,j_wb = gradient_descent(x_train, y_train, alpha, w_init,b_init, num_iters, costfn, gradientfn)

print(f'w_final:{w_final:.2f}, b_final:{b_final:.2f}')

print(j_wb)

```