

## Week 2

### Multiple features

#### Multiple features (variables)

	Size in feet <sup>2</sup>	Number of bedrooms	Number of floors	Age of home in years	Price (\$) in \$1000's
	$x_1$	$x_2$	$x_3$	$x_4$	
$i=2$	2104	5	1	45	460
	1416	3	(2)	40	232
	1534	3	2	30	315
	852	2	1	36	178
	...	...	...	...	...

$x_j = j^{\text{th}}$  feature  
 $n$  = number of features  
 $\vec{x}^{(i)}$  = features of  $i^{\text{th}}$  training example  
 $x_j^{(i)}$  = value of feature  $j$  in  $i^{\text{th}}$  training example

$j=1 \dots 4$   
 $n=4$   
 $\vec{x}^{(2)} = [1416 \ 3 \ (2) \ 40]$   
 $x_3^{(2)} = 2$

$$f_{\vec{w},b}(\vec{x}) = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$$

$\vec{w} = [w_1 \ w_2 \ w_3 \ \dots \ w_n]$  parameters of the model  
 $b$  is a number  
 $\vec{x} = [x_1 \ x_2 \ x_3 \ \dots \ x_n]$  vector

$$f_{\vec{w},b}(\vec{x}) = \vec{w} \cdot \vec{x} + b = w_1x_1 + w_2x_2 + w_3x_3 + \dots + w_nx_n + b$$

dot product      multiple linear regression  
 (not multivariate regression)

### Vectorization part 1

Parameters and features

$$\vec{w} = [w_1 \ w_2 \ w_3] \quad n=3$$

$b$  is a number

$$\vec{x} = [x_1 \ x_2 \ x_3]$$

linear algebra: count from 1

NumPy

$w = \text{np.array}([1.0, 2.5, -3.3])$

$b = 4$

$x = \text{np.array}([10, 20, 30])$

code: count from 0

Without vectorization  $n=100,000$

$$f_{\vec{w},b}(\vec{x}) = w_1x_1 + w_2x_2 + w_3x_3 + b$$

```
f = w[0] * x[0] +
    w[1] * x[1] +
    w[2] * x[2] + b
```



Without vectorization

$$f_{\vec{w},b}(\vec{x}) = \left( \sum_{j=1}^n w_j x_j \right) + b \quad \sum_{j=1}^n \rightarrow j=1 \dots n$$

range(0, n) → j = 0, ..., n-1

```
f = 0
for j in range(0, n):
    f = f + w[j] * x[j]
f = f + b
```



Vectorization

$$f_{\vec{w},b}(\vec{x}) = \vec{w} \cdot \vec{x} + b$$

```
f = np.dot(w, x) + b
```

### Vectorization part 2

## Without vectorization

```
for j in range(0,16):
    f = f + w[j] * x[j]
```

$$t_0 \quad f + w[0] * x[0]$$

$$t_1 \quad f + w[1] * x[1]$$

$$\dots$$

$$t_{15} \quad f + w[15] * x[15]$$

## Vectorization

```
np.dot(w,x)
```

$t_0$

w[0] w[1] ... w[15]

in parallel \*

x[0] x[1] ... x[15]

$t_1$

w[0]\*x[0] + w[1]\*x[1] + ... + w[15]\*x[15]

efficient → scale to large datasets

## Gradient descent

$$\vec{w} = (w_1 \ w_2 \ \dots \ w_{16}) \quad \text{parameters}$$

$$\text{derivatives } \vec{d} = (d_1 \ d_2 \ \dots \ d_{16})$$

```
w = np.array([0.5, 1.3, ... 3.4])
```

```
d = np.array([0.3, 0.2, ... 0.4])
```

compute  $w_j = w_j - 0.1d_j$  for  $j = 1 \dots 16$

### Without vectorization

$$w_1 = w_1 - 0.1d_1$$

$$w_2 = w_2 - 0.1d_2$$


⋮

$$w_{16} = w_{16} - 0.1d_{16}$$

```
for j in range(0,16):
    w[j] = w[j] - 0.1 * d[j]
```

### With vectorization

$$\vec{w} = \vec{w} - 0.1\vec{d}$$



```
w = w - 0.1 * d
```

remember here, w is a vector containing multiple values of w's

## Gradient descent for multiple linear regression

### Previous notation

Parameters  $w_1, \dots, w_n$   
 $b$

Model  $f_{\vec{w},b}(\vec{x}) = w_1x_1 + \dots + w_nx_n + b$

Cost function  $J(w_1, \dots, w_n, b)$

### Vector notation

$\vec{w} = [w_1 \ \dots \ w_n]$  ← vector of length n

$b$  still a number

$f_{\vec{w},b}(\vec{x}) = \vec{w} \cdot \vec{x} + b$

$J(\vec{w}, b)$  dot product

## Gradient descent

repeat {

$$w_j = w_j - \alpha \frac{\partial}{\partial w_j} J(w_1, \dots, w_n, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(w_1, \dots, w_n, b)$$

}

repeat {

$$w_j = w_j - \alpha \frac{\partial}{\partial w_j} J(\vec{w}, b)$$

$$b = b - \alpha \frac{\partial}{\partial b} J(\vec{w}, b)$$

}

## Gradient descent

### One feature

repeat {

$$w = w - \alpha \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)}) x^{(i)}$$

$$\rightarrow \frac{\partial}{\partial w} J(w, b)$$

$$b = b - \alpha \frac{1}{m} \sum_{i=1}^m (f_{w,b}(x^{(i)}) - y^{(i)})$$

simultaneously update  $w, b$

}

### n features ( $n \geq 2$ )

repeat {

$$w_1 = w_1 - \alpha \frac{1}{m} \sum_{i=1}^m (f_{\vec{w},b}(\vec{x}^{(i)}) - y^{(i)}) x_1^{(i)}$$

⋮

$$w_n = w_n - \alpha \frac{1}{m} \sum_{i=1}^m (f_{\vec{w},b}(\vec{x}^{(i)}) - y^{(i)}) x_n^{(i)}$$

$$b = b - \alpha \frac{1}{m} \sum_{i=1}^m (f_{\vec{w},b}(\vec{x}^{(i)}) - y^{(i)})$$

simultaneously update  
 $w_j$  (for  $j = 1, \dots, n$ ) and  $b$

}

General Notation	Description	Python (if applicable)
$a$	scalar, non bold	
$\mathbf{a}$	vector, bold	
$\mathbf{A}$	matrix, bold capital	
<b>Regression</b>		
$\mathbf{X}$	training example matrix	<code>X_train</code>
$\mathbf{y}$	training example targets	<code>y_train</code>
$\mathbf{x}^{(i)}, y^{(i)}$	$i$ th Training Example	<code>X[i] , y[i]</code>
$m$	number of training examples	<code>m</code>
$n$	number of features in each example	<code>n</code>
$\mathbf{w}$	parameter: weight,	<code>w</code>
$b$	parameter: bias	<code>b</code>
$f_{\mathbf{w},b}(\mathbf{x}^{(i)})$	The result of the model evaluation at $\mathbf{x}^{(i)}$ parameterized by $\mathbf{w}, b$ : $f_{\mathbf{w},b}(\mathbf{x}^{(i)}) = \mathbf{w} \cdot \mathbf{x}^{(i)} + b$	<code>f_wb</code>

```
X_train = np.array([[2104, 5, 1, 45], [1416, 3, 2, 40], [852, 2, 1, 35]])
y_train = np.array([460, 232, 178])
```

## 2.1 Matrix X containing our examples

Similar to the table above, examples are stored in a NumPy matrix `X_train`. Each row of the matrix represents one example. When you have  $m$  training examples ( $m$  is three in our example), and there are  $n$  features (four in our example),  $\mathbf{X}$  is a matrix with dimensions  $(m, n)$  ( $m$  rows,  $n$  columns).

$$\mathbf{X} = \begin{pmatrix} x_0^{(0)} & x_1^{(0)} & \cdots & x_{n-1}^{(0)} \\ x_0^{(1)} & x_1^{(1)} & \cdots & x_{n-1}^{(1)} \\ \cdots & \cdots & \cdots & \cdots \\ x_0^{(m-1)} & x_1^{(m-1)} & \cdots & x_{n-1}^{(m-1)} \end{pmatrix}$$

notation:

- $\mathbf{x}^{(i)}$  is vector containing example  $i$ .  $\mathbf{x}^{(i)} = (x_0^{(i)}, x_1^{(i)}, \dots, x_{n-1}^{(i)})$
- $x_j^{(i)}$  is element  $j$  in example  $i$ . The superscript in parenthesis indicates the example number while the subscript represents an element.

Display the input data.

```
import copy, math
import numpy as np
import matplotlib.pyplot as plt

X_train = np.array([[2104, 5, 1, 45], [1416, 3, 2, 40], [852, 2, 1, 35]])
y_train = np.array([460, 232, 178])

# data is stored in numpy array/matrix
print(f"X Shape: {X_train.shape}, X Type:{type(X_train)}")
print(X_train)
print(f"y Shape: {y_train.shape}, y Type:{type(y_train)}")
print(y_train)
"""X Shape: (3, 4), X Type:<class 'numpy.ndarray'>
[[2104  5  1  45]
 [1416  3  2  40]
 [ 852  2  1  35]]
y Shape: (3,), y Type:<class 'numpy.ndarray'>
[460 232 178]"""
```

- $\mathbf{w}$  is a vector with  $n$  elements.
  - Each element contains the parameter associated with one feature.
  - in our dataset,  $n$  is 4.
  - notionally, we draw this as a column vector

$$\mathbf{w} = \begin{pmatrix} w_0 \\ w_1 \\ \cdots \\ w_{n-1} \end{pmatrix}$$

- $b$  is a scalar parameter.

```
b_init = 785.1811367994083
w_init = np.array([ 0.39133535, 18.75376741, -53.36032453, -26.42131618]) #w and b are selected for optimal solns as they are better
print(f"w_init shape: {w_init.shape}, b_init type: {type(b_init)}")
```

$$f_{\mathbf{w},b}(\mathbf{x}) = w_0x_0 + w_1x_1 + \dots + w_{n-1}x_{n-1} + b$$

$$f_{\mathbf{w},b}(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x} + b$$

## Single Prediction element by element

```
def predict_single_loop(x, w, b):
    n = x.shape[0]
    p = 0
    for i in range(n):
        p_i = x[i] * w[i]
        p = p + p_i
    p = p + b
    return p
# row wise
'''

# get a row from our training data
x_vec = X_train[0,:]
print(f"x_vec shape {x_vec.shape}, x_vec value: {x_vec}")

# make a prediction
f_wb = predict_single_loop(x_vec, w_init, b_init)
print(f"f_wb shape {f_wb.shape}, prediction: {f_wb}")

x_vec shape (4,), x_vec value: [2104    5    1   45]
f_wb shape (), prediction: 459.9999976194083 '''
```

## Single Prediction, vector

```
def predict(x, w, b):
    p = np.dot(x, w) + b
    return p
```

## Compute Cost With Multiple Variables

$$J(\mathbf{w}, b) = \frac{1}{2m} \sum_{i=0}^{m-1} (f_{\mathbf{w},b}(\mathbf{x}^{(i)}) - y^{(i)})^2$$

$$f_{\mathbf{w},b}(\mathbf{x}^{(i)}) = \mathbf{w} \cdot \mathbf{x}^{(i)} + b$$

```
def compute_cost(X, y, w, b):
    m = X.shape[0]
    cost = 0.0
    for i in range(m):
        f_wb_i = np.dot(X[i], w) + b          #(n,)(n,) = scalar (see np.dot)
        cost = cost + (f_wb_i - y[i])**2      #scalar
    cost = cost / (2 * m)                     #scalar
    return cost
```

## Gradient Descent with multiple variables

Gradient descent for multiple variables:

repeat until convergence: {

$$w_j = w_j - \alpha \frac{\partial J(\mathbf{w}, b)}{\partial w_j} \quad \text{for } j = 0..n-1$$
$$b = b - \alpha \frac{\partial J(\mathbf{w}, b)}{\partial b}$$

}

where, n is the number of features, parameters  $w_j$ ,  $b$ , are updated simultaneously and where

$$\frac{\partial J(\mathbf{w}, b)}{\partial w_j} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{\mathbf{w},b}(\mathbf{x}^{(i)}) - y^{(i)}) x_j^{(i)}$$
$$\frac{\partial J(\mathbf{w}, b)}{\partial b} = \frac{1}{m} \sum_{i=0}^{m-1} (f_{\mathbf{w},b}(\mathbf{x}^{(i)}) - y^{(i)})$$

- m is the number of training examples in the data set
- $f_{\mathbf{w},b}(\mathbf{x}^{(i)})$  is the model's prediction, while  $y^{(i)}$  is the target value

```
def compute_gradient(X, y, w, b):
    m, n = X.shape          #(number of examples, number of features)
    dj_dw = np.zeros((n,))
    dj_db = 0.

    for i in range(m):
        err = (np.dot(X[i], w) + b) - y[i]
        for j in range(n):
```

```

        dj_dw[j] = dj_dw[j] + err * X[i, j]
    dj_db = dj_db + err
    dj_dw = dj_dw / m
    dj_db = dj_db / m

    return dj_db, dj_dw

```

```

#Compute and display gradient
tmp_dj_db, tmp_dj_dw = compute_gradient(X_train, y_train, w_init, b_init)
print(f'dj_db at initial w,b: {tmp_dj_db}')
print(f'dj_dw at initial w,b: \n {tmp_dj_dw}')

dj_db at initial w,b: -1.673925169143331e-06
dj_dw at initial w,b:
[-2.73e-03 -6.27e-06 -2.22e-06 -6.92e-05]

```

## Gradient Descent With Multiple Variables

```

def gradient_descent(X, y, w_in, b_in, cost_function, gradient_function, alpha, num_iters):

    # An array to store cost J and w's at each iteration primarily for graphing later
    J_history = []
    w = copy.deepcopy(w_in) #avoid modifying global w within function
    b = b_in

    for i in range(num_iters):

        # Calculate the gradient and update the parameters
        dj_db,dj_dw = gradient_function(X, y, w, b) ##None

        # Update Parameters using w, b, alpha and gradient
        w = w - alpha * dj_dw ##None
        b = b - alpha * dj_db ##None

        '''# Save cost J at each iteration
        if i<100000: # prevent resource exhaustion
            J_history.append( cost_function(X, y, w, b))

        # Print cost every at intervals 10 times or as many iterations if < 10
        if i% math.ceil(num_iters / 10) == 0:
            print(f"Iteration {i:4d}: Cost {J_history[-1]:8.2f}  ")'''

    return w, b, J_history #return final w,b and J history for graphing

#-----client code
# initialize parameters
initial_w = np.zeros_like(w_init)
initial_b = 0.
# some gradient descent settings
iterations = 1000
alpha = 5.0e-7
# run gradient descent
w_final, b_final, J_hist = gradient_descent(X_train, y_train, initial_w, initial_b,
                                             compute_cost, compute_gradient,
                                             alpha, iterations)

print(f"b,w found by gradient descent: {b_final:0.2f},{w_final} ")
m,_ = X_train.shape
for i in range(m):
    print(f"prediction: {np.dot(X_train[i], w_final) + b_final:0.2f}, target value: {y_train[i]}")

```

## final code using vectorization

```

import numpy as np

import copy

'''
x_train = np.array([1.0, 2.0])

y_train = np.array([300.0, 500.0])

w_init = 0

b_init = 0

alpha = 0.001

num_iters = 100000

```

```

def costfn(x_train, y_train, w,b):

    j_wb = 0

    m = x_train.shape[0]

    for i in range(m):

        f_wb = w*x_train[i]+b

        j_wb = j_wb + (f_wb-y_train[i])**2

    j_wb = j_wb/(2*m)

    return j_wb

def gradientfn(x_train, y_train, w,b):

    m = x_train.shape[0]

    d_jwb_w = 0

    d_jwb_b = 0

    for i in range(m):

        f_wb = w*x_train[i] + b

        d_jwb_w = d_jwb_w + (f_wb-y_train[i])*x_train[i]

        d_jwb_b = d_jwb_b + (f_wb-y_train[i])

    d_jwb_w = d_jwb_w/m

    d_jwb_b = d_jwb_b/m

    return d_jwb_w, d_jwb_b

def gradient_descent(x_train, y_train, alpha, w_init, b_init, num_iters, costfn, gradientfn):

    w,b = w_init, b_init

    dj_dw , dj_db = gradientfn(x_train, y_train, w,b)

    jwb = []

    for i in range(num_iters):

        dj_dw , dj_db = gradientfn(x_train, y_train, w,b)

        w = w - alpha*dj_dw

        b = b - alpha*dj_db

        if (i%1000)==0:

            jwb.append(costfn(x_train, y_train, w,b))

    return w,b,jwb

w_final, b_final,j_wb = gradient_descent(x_train, y_train, alpha, w_init,b_init, num_iters, costfn, gradientfn)

print(f'w_final:{w_final:.2f}, b_final:{b_final:.2f}')

print(j_wb)
'''

X_train = np.array([[2104, 5, 1, 45], [1416, 3, 2, 40], [852, 2, 1, 35]])

```

```

y_train = np.array([460, 232, 178])

alpha = 5.0e-7

num_iters = 1000


b_init = 785.1811367994083

w_init = np.array([ 0.39133535, 18.75376741, -53.36032453, -26.42131618])


'''def predict_single_loop(x, y, w, b):

    n = x.shape[0]

    f_wb = 0

    for i in range(n):

        f_wb += w[i]*x[i]

    f_wb +=b

    return f_wb

print(predict_single_loop(X_train[0], y_train, w_init, b_init))'''

'''

def predict(x,w,b):

    f_wb = np.dot(x,w) + b

    print(f_wb)

    #here f_wb is a list of predictions, number of f_wb is equal to x.shape[0] (the number of rows)

predict(X_train, w_init,b_init)

'''


def compute_cost(X, y, w,b):

    m = X.shape[0]

    cost = 0

    for i in range(m):

        f_wb_i = np.dot(X[i],w)+b

        cost += (f_wb_i-y[i])**2

    cost = cost/(2*m)

    return cost


def compute_gradient(X, y, w,b):

    m,n = X.shape

    dj_dw = np.zeros(n)

```

```

dj_db = 0

for i in range(m):

    error = np.dot(X[i],w)+b - y[i]

    for j in range(n):

        dj_dw[j] += error*X[i,j]

    dj_db += error

dj_db /= m

dj_dw /= m


return dj_dw, dj_db


def gradient_descent(X_train, y_train, alpha, num_iters, w_init, b_init, compute_cost, compute_gradient):

    jwb = []

    w = copy.deepcopy(w_init)

    b = b_init

    for i in range(num_iters):

        dj_dw, dj_db = compute_gradient(X_train, y_train, w, b)

        w = w - alpha*dj_dw

        b = b - alpha*dj_db

        if i%100==0:

            jwb.append(compute_cost(X_train, y_train, w_init, b_init))

    return w,b,jwb


b_init = 0

w_init = np.zeros(X_train.shape[1])


w_final,b_final,j_wb = gradient_descent(X_train, y_train, alpha, num_iters, w_init, b_init, compute_cost, compute_gradient)

print(w_final, b_final)


print(j_wb)


f_wb = np.dot(X_train,w_final) + b_final

print(f_wb)

```

## Feature scaling Part 1



## Feature and parameter values

$$\widehat{\text{price}} = w_1 x_1 + w_2 x_2 + b$$

$\downarrow$   
size
 $\downarrow$   
# bedrooms
 $\downarrow$   
large
 $\downarrow$   
small

$x_1$ : size (feet<sup>2</sup>)  
 range: 300 – 2,000  
 $x_2$ : # bedrooms  
 range: 0 – 5

House:  $x_1 = 2000$ ,  $x_2 = 5$ ,  $\text{price} = \$500\text{k}$  one training example

size of the parameters  $w_1, w_2$ ?

$$w_1 = 50, w_2 = 0.1, b = 50$$

$$\widehat{\text{price}} = \underbrace{50 * 2000}_{100,000\text{K}} + \underbrace{0.1 * 5}_{0.5\text{K}} + \underbrace{50}_{50\text{K}}$$





$$\widehat{\text{price}} = \$100,050.5\text{k} = \$100,050,500$$

$$w_1 = 0.1, w_2 = 50, b = 50$$


$$\widehat{\text{price}} = \underbrace{0.1 * 2000\text{k}}_{200\text{K}} + \underbrace{50 * 5}_{250\text{K}} + \underbrace{50}_{50\text{K}}$$

$$\widehat{\text{price}} = \$500\text{k} \text{ more reasonable}$$


## Feature size and parameter size

	size of feature $x_j$	size of parameter $w_j$
size in feet <sup>2</sup>		
# bedrooms		

### Features




### Parameters




## Feature size and gradient descent

### Features



### Parameters



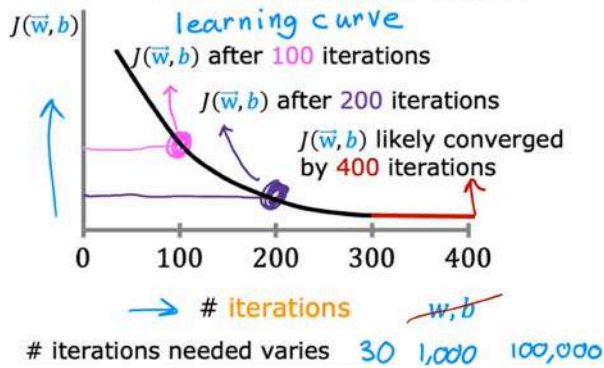
the cost function bounces a lot until it reaches the center for a feature set not scaled.

## Feature scaling part 2 - in notes

## Checking gradient descent for convergence

# Make sure gradient descent is working correctly

objective:  $\min_{\bar{w}, b} J(\bar{w}, b)$   $J(\bar{w}, b)$  should **decrease** after every iteration




Automatic convergence test

Let  $\epsilon$  "epsilon" be  $10^{-3}$ .  
 $0.001$


If  $J(\bar{w}, b)$  decreases by  $\leq \epsilon$  in one iteration, declare **convergence**.

(found parameters  $\bar{w}, b$  to get close to global minimum)

Identify problem with gradient descent




Adjust learning rate



Values of  $\alpha$  to try:

... 0.001 0.003 0.01 0.03 0.1 0.3 1 ...  
3x  $\approx 3x$  3x  $\approx 3x$  3x  $\approx 3x$



how to choose alpha effectively

## C1\_W2\_Lab03\_Feature\_Scaling\_and\_Learning\_Rate\_Solne

- After z-score normalization, all features will have a mean of 0 and a standard deviation of 1.

$$x_j^{(i)} = \frac{x_j^{(i)} - \mu_j}{\sigma_j}$$

- where  $j$  selects a feature or a column in the  $XX$  matrix.  $\mu_j$  is the mean of all the values for feature ( $j$ ) and  $\sigma_j$  is the standard deviation of feature ( $j$ ).

$$\mu_j = \frac{1}{m} \sum_{i=0}^{m-1} x_j^{(i)}$$

$$\sigma_j^2 = \frac{1}{m} \sum_{i=0}^{m-1} (x_j^{(i)} - \mu_j)^2$$

- Implementation Note:** When normalizing the features, it is important to store the values used for normalization - the mean value and the standard deviation used for the computations. After learning the parameters from the model, we often want to predict the prices of houses

we have not seen before. Given a new x value (living room area and number of bedrooms), we must first normalize x using the mean and standard deviation that we had previously computed from the training set.

```
def zscore_normalize_features(X):
    # find the mean of each column/feature
    mu = np.mean(X, axis=0) # mu will have shape (n,)
    # find the standard deviation of each column/feature
    sigma = np.std(X, axis=0) # sigma will have shape (n,)
    # element-wise, subtract mu for that column from each example, divide by std for that column
    X_norm = (X - mu) / sigma

    return (X_norm, mu, sigma)

"""or just"""

mu = np.mean(X_train,axis=0)
sigma = np.std(X_train,axis=0)
X_mean = (X_train - mu)
X_norm = (X_train - mu)/sigma
```

A few points to note:

- with multiple features, we can no longer have a single plot showing results versus features.
- when generating the plot, the normalized features were used. Any predictions using the parameters learned from a normalized training set must also be normalized.

```
"""**Prediction** The point of generating our model is to use it to predict housing prices that are not in the data set. Let's
predict the price of a house with 1200 sqft, 3 bedrooms, 1 floor, 40 years old. Recall, that you must normalize the data with the
mean and standard deviation derived when the training data was normalized."""
# First, normalize out example.
x_house = np.array([1200, 3, 1, 40])
x_house_norm = (x_house - X_mu) / X_sigma
print(x_house_norm)
x_house_predict = np.dot(x_house_norm, w_norm) + b_norm
print(f" predicted price of a house with 1200 sqft, 3 bedrooms, 1 floor, 40 years old = ${x_house_predict*1000:0.0f}")

#output
#[-0.53  0.43 -0.79  0.06]
#predicted price of a house with 1200 sqft, 3 bedrooms, 1 floor, 40 years old = $318709
```

## Feature engineering

### Feature engineering

$$f_{\underline{w},b}(\underline{x}) = \underbrace{w_1 x_1}_{\text{frontage}} + \underbrace{w_2 x_2}_{\text{depth}} + b$$

$$\text{area} = \text{frontage} \times \text{depth}$$

$$x_3 = x_1 x_2$$

new feature


$$f_{\underline{w},b}(\underline{x}) = \underbrace{w_1 x_1}_{\text{frontage}} + \underbrace{w_2 x_2}_{\text{depth}} + \underbrace{w_3 x_3}_{\text{area}} + b$$



Feature engineering:  
Using intuition to design  
new features, by  
transforming or combining  
original features.

here we don't just remove the features we used to make the new feature but here we add it so that the model can see how useful or not the 3 features are, let the model decide, just let all of them be there.

## polynomial regression



## C1\_W2\_Lab04\_FeatEng\_PolyReg\_Soln

```
x = np.array([1, 2, 3, 4])
X = x.reshape(-1, 1)
will give you:
[[1],
 [2],
 [3],
 [4]]
### Breakdown:
- '-1' means "figure out the right number of rows automatically."
- '1' means "make it 1 column."
- This format ('n x 1') is especially useful when feeding data into machine learning models like 'sklearn', which expect a 2D array of features.
```

```
x = np.arange(0, 20, 1)
will give you:
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9,
        10, 11, 12, 13, 14, 15, 16, 17, 18, 19])
0-start, 20-stop, 1-step
```

in feature engineering, we may have an intuition about what features may be useful or not, giving those features to GD will help us know if they impact the end result more or less and whether they are imp or not. basically GD will tell us whether the new features we engineered are good or bad but giving is a low or high weight.

it is important to do normalization when applying feature engineering. eg. x, x2, x3 need to be normalized as they would have varying scale


```
In [11]: x = np.arange(0,20,1)
y = np.cos(x/2)

X = np.c_[x, x**2, x**3, x**4, x**5, x**6, x**7, x**8, x**9, x**10, x**11, x**12, x**13]
X = zscore_normalize_features(X)

model_w, model_b = run_gradient_descent_feng(X, y, iterations=1000000, alpha = 1e-1)

plt.scatter(x, y, marker='x', c='r', label="Actual Value"); plt.title("Normalized x x**2, x**3 feature")
plt.plot(x, X@model_w + model_b, label="Predicted Value"); plt.xlabel("x"); plt.ylabel("y"); plt.legend(); plt.show()
```

Iteration 0, Cost: 2.20188e-01  
 Iteration 100000, Cost: 1.70074e-02  
 Iteration 200000, Cost: 1.27603e-02  
 Iteration 300000, Cost: 9.73032e-03  
 Iteration 400000, Cost: 7.56440e-03  
 Iteration 500000, Cost: 6.01412e-03  
 Iteration 600000, Cost: 4.90251e-03  
 Iteration 700000, Cost: 4.10351e-03  
 Iteration 800000, Cost: 3.52730e-03  
 Iteration 900000, Cost: 3.10989e-03  
 w,b found by gradient descent: w: [ -1.34 -10. 24.78 5.96 -12.49 -16.26 -9.51 0.59 8.7 11.94  
 9.27 0.79 -12.82], b: -0.0073



## C1\_W2\_Lab05\_Sklearn\_GD\_Soln

```

"""code to implement linear regression using GD using scikit-learn"""
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import SGDRegressor
from sklearn.preprocessing import StandardScaler
from lab_utils_multi import load_house_data
from lab_utils_common import dlc
np.set_printoptions(precision=2)
plt.style.use('./deeplearning.mplstyle')

"""
- Scikit-learn has a gradient descent regression model [sklearn.linear_model.SGDRegressor] Like your previous implementation of
gradient descent, this model performs best with normalized inputs.
- [sklearn.preprocessing.StandardScaler] will perform z-score normalization as in a previous lab. Here it is referred to as
'standard score'.
"""

X_train, y_train = load_house_data()
X_features = ['size(sqft)', 'bedrooms', 'floors', 'age']
scaler = StandardScaler()
X_norm = scaler.fit_transform(X_train)
print(f"Peak to Peak range by column in Raw X:{np.ptp(X_train,axis=0)}")
print(f"Peak to Peak range by column in Normalized X:{np.ptp(X_norm,axis=0)}")

sgdr = SGDRegressor(max_iter=1000)
sgdr.fit(X_norm, y_train)
print(sgdr)
print(f"number of iterations completed: {sgdr.n_iter_}, number of weight updates: {sgdr.t_}")

b_norm = sgdr.intercept_
w_norm = sgdr.coef_
print(f"model parameters: w: {w_norm}, b:{b_norm}")
# both below give same output, 2 diff ways of doing
# make a prediction using sgdr.predict()
y_pred_sgd = sgdr.predict(X_norm)
# make a prediction using w,b.
y_pred = np.dot(X_norm, w_norm) + b_norm

```

check documentation of scikit here - [scikit-learn: machine learning in Python — scikit-learn 1.6.1 documentation](https://scikit-learn.org/stable/user_guide.html) user guide or api

## C1\_W2\_Linear\_Regression

1-D numpy array's shape gives - (m,)

2-D numpy array's shape gives - (m,n)

### Visualize your data

It is often useful to understand the data by visualizing it.

- For this dataset, you can use a scatter plot to visualize the data, since it has only two properties to plot (profit and population).
- Many other problems that you will encounter in real life have more than two properties (for example, population, average household income, monthly profits, monthly sales). When you have more than two properties, you can still use a scatter plot to see the relationship between each pair of properties.